



Outlines

Overview
Introduction
Linear Algebra
Probability
Linear Regression 1
Linear Regression 2
Linear Classification 1
Linear Classification 2
Kernel Methods
Sparse Kernel Methods
Mixture Models and EM 1
Mixture Models and EM 2
Neural Networks 1
Neural Networks 2
Principal Component Analysis
Autoencoders
Graphical Models 1
Graphical Models 2
Graphical Models 3
Sampling
Sequential Data 1
Sequential Data 2

Statistical Machine Learning

Assignment Project Exam Help

Christian Walder

Machine Learning Research Group

CSIRO Data61

and

College of Engineering and Computer Science

The Australian National University

Canberra

Semester One, 2020.

<https://powcoder.com>

Add WeChat powcoder

(Many figures from C. M. Bishop, "Pattern Recognition and Machine Learning")



Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

Assignment Project Exam Help

- Estimate best predictor = training = learning

Given data $(x_1, y_1), \dots, (x_n, y_n)$, find a predictor $f_{\mathbf{w}}(\cdot)$.

- 1 Identify the type of input x and output y data
- 2 Propose a (linear) mathematical model for $f_{\mathbf{w}}$
- 3 Design an objective function or likelihood
- 4 Calculate the optimal parameter (\mathbf{w})
- 5 Model uncertainty using the Bayesian approach
- 6 Implement and compute (the algorithm in python)
- 7 Interpret and diagnose results

<https://powcoder.com>

Add WeChat powcoder



Classification



Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

- Goal : Given input data \mathbf{x} , assign it to one of K discrete classes \mathcal{C}_k where $k = 1, \dots, K$.
- Divide the input space into different regions.
- Equivalently: map each point to a categorical label.



Length of petal [in cm] vs sepal [cm] for three types of flowers
(Iris Setosa, Iris Versicolor, Iris Virginica).

How to represent binary class labels?



Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

Assignment Project Exam Help

- Class labels are no longer real values as in regression, but a discrete set.
- Two classes : $t \in \{0, 1\}$
($t = 1$ represents class C_1 and $t = 0$ represents class C_0)
- Can interpret the value of t as the probability of class C_1 , with only two values possible for the probability, 0 or 1.
- Note: Other conventions to map classes into integers possible, check the setup.

<https://powcoder.com>

Add WeChat powcoder

How to represent multi-class labels?



Assignment Project Exam Help

- If there are more than two classes ($K > 2$), we call it a multi-class setup.
- Often used: 1-of- K coding scheme in which \mathbf{t} is a vector of length K which has all values 0 except for $t_j = 1$, where j comes from the membership in class C_j to encode.
- Example: Given 5 classes, $\{C_1, \dots, C_5\}$. Membership in class C_2 will be encoded as the target vector

$\mathbf{t} = (0, 1, 0, 0, 0)^T$

Add WeChat powcoder

- Note: Other conventions to map multi-classes into integers possible, check the setup.



Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

- Idea: Use again a Linear Model as in regression: $y(\mathbf{x}, \mathbf{w})$ is a linear function of the parameters \mathbf{w}

$$y(\mathbf{x}_n, \mathbf{w}) = \mathbf{w}^\top \phi(\mathbf{x}_n)$$

- But generally $y(\mathbf{x}_n, \mathbf{w}) \in \mathbb{R}$.
Example: Which class is $y(\mathbf{x}, \mathbf{w}) = 0.71623$?





- Apply a mapping $f: \mathbb{R} \rightarrow \mathbb{Z}$ to the linear model to get the discrete class labels.

Assignment Project Exam Help

$$y(\mathbf{x}_n, \mathbf{w}) = f(\mathbf{w}^\top \phi(\mathbf{x}_n))$$

- Activation function: $f(\cdot)$
- Link function: $f^{-1}(\cdot)$

<https://powcoder.com>

Add WeChat powcoder




Figure: Example of an activation function $f(z) = \text{sign}(z)$.

Three Models for Decision Problems



Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

- Find a **discriminant function** $f(\mathbf{x})$ which maps each input directly onto a class label

- Discriminative Models

- 1 Solve the inference problem of determining the posterior class probabilities $p(C_k | \mathbf{x})$.
- 2 Use decision theory to assign each new \mathbf{x} to one of the classes

- Generative Models

- 1 Solve the inference problem of determining the class conditional probabilities $p(\mathbf{x} | C_k)$.
- 2 Also, infer the prior class probabilities $p(C_k)$.
- 3 Use Bayes' theorem to find the posterior $p(C_k | \mathbf{x})$.
- 4 Alternatively, model the joint distribution $p(\mathbf{x}, C_k)$ directly.
- 5 Use decision theory to assign each new \mathbf{x} to one of the classes.



Definition

A **discriminant** is a function that maps from an input vector \mathbf{x} to one of K classes, denoted by \mathcal{C}_k .

- Consider first two classes ($K = 2$).
- Construct a linear function of the inputs \mathbf{x}

$$y(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + w_0$$

such that \mathbf{x} being assigned to class \mathcal{C}_1 if $y(\mathbf{x}) \geq 0$, and to class \mathcal{C}_2 otherwise.

- **weight vector** \mathbf{w}
- **bias** w_0 (sometimes $-w_0$ called **threshold**)






Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm



Assignment Project Exam Help

<https://powcoder.com>

- Gradient = direction of steepest ascent = $\nabla_x w^T x = w$.
- The set $w^T x + w_0 = 0$ is a hyper-plane.

- Projecting x on that hyper-plane means finding

$\arg \min_x \|x - x_\perp\|$ subject to the constraint $w^T x_\perp + w_0 = 0$. Geometrically: move in the direction $\frac{w}{\|w\|}$.

- Rate of change of function value in that direction is

$$\frac{d}{da} \left(a \frac{w}{\|w\|} \right)^T w = a \|w\|.$$

- The length $\left\| a \frac{w}{\|w\|} \right\| = \frac{a}{\|w\|} \|w\| = a$.

- For a fixed change in $w^T \left(a \frac{w}{\|w\|} \right)$, $a \propto \frac{1}{\|w\|}$.

Add WeChat in the powcoder



Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

Assignment Project Exam Help

- Decision boundary $y(\mathbf{x}) = 0$ is a $(D - 1)$ -dimensional hyperplane in a D -dimensional input space (decision surface).

- \mathbf{w} is orthogonal to any vector lying in the decision surface.

- Proof: Assume \mathbf{x}_A and \mathbf{x}_B are two points lying in the decision surface. Then,

$0 = y(\mathbf{x}_A) = y(\mathbf{x}_B) \Rightarrow \mathbf{w}^T (\mathbf{x}_A - \mathbf{x}_B)$

Two Classes




- $y(\mathbf{x})$ gives a **signed** measure of the perpendicular distance r **from** the decision surface **to** \mathbf{x} , that is $r = y(\mathbf{x})/\|\mathbf{w}\|$.

$$y(\mathbf{x}) = \mathbf{w}^T \left(\mathbf{x}_\perp + r \frac{\mathbf{w}}{\|\mathbf{w}\|} \right) + w_0 = r \frac{\mathbf{w}^T \mathbf{w}}{\|\mathbf{w}\|} + \overbrace{\mathbf{w}^T \mathbf{x}_\perp}^0 + w_0 = r \|\mathbf{w}\|$$

<https://powcoder.com>

Add WeChat powcoder






- The normal distance from the origin to the decision surface is therefore

$$-\frac{y(\mathbf{0})}{\|\mathbf{w}\|} = -\frac{w_0}{\|\mathbf{w}\|}$$

<https://powcoder.com>

Add WeChat powcoder





Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

Assignment Project Exam Help

- More compact notation : Add an extra dimension to the input space and set the value to $x_0 = 1$.
- Also define $\tilde{\mathbf{w}} = (w_0, \mathbf{w})$ and $\tilde{\mathbf{x}} = (1, \mathbf{x})$

<https://powcoder.com>

$$y(\mathbf{x}) = \tilde{\mathbf{w}}^\top \tilde{\mathbf{x}}$$

(if it helps, you may think of $\tilde{\mathbf{w}}^\top$ as a function).

- Decision surface is now a D -dimensional hyperplane in a $D + 1$ -dimensional expanded input space.

Add WeChat powcoder




- Number of classes $K > 2$
- Can we combine a number of two-class discriminant functions using $K - 1$ one-versus-the-rest classifiers?

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder






- Number of classes $K > 2$
- Can we combine a number of two-class discriminant functions using $K(K-1)/2$ one-versus-one classifiers?

<https://powcoder.com>

Add WeChat powcoder





Classification

Generalised Linear
Model

Discriminant Functions


Fisher's Linear
Discriminant

The Perceptron
Algorithm

- Number of classes $K > 2$
- Solution: Use K linear functions

$$y_k(\mathbf{x}) = \mathbf{w}_k^\top \mathbf{x} + w_{k0}$$

- Assign input \mathbf{x} to class C_k if $y_k(\mathbf{x}) > y_j(\mathbf{x})$ for all $j \neq k$.
- Decision boundary between class C_k and C_j given by $y_k(\mathbf{x}) = y_j(\mathbf{x})$



Add WeChat powcoder



Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

Assignment Project Exam Help

- Regression with a linear function of the model parameters and minimisation of sum-of-squares error function resulted in a closed-form solution for the parameter values.

- Is this also possible for classification?
- Given input data \mathbf{x} belonging to one of K classes C_k .
- Use 1-of- K binary coding scheme.

- Each class is described by its own linear model

$$y_k(\mathbf{x}) = \mathbf{w}_k^\top \mathbf{x} + w_{k0} \quad k = 1, \dots, K$$

<https://powcoder.com>
Add WeChat powcoder



Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

- With the conventions

$$\tilde{\mathbf{w}}_k = \begin{bmatrix} \mathbf{w}_k \\ w_k \end{bmatrix}$$

$$\in \mathbb{R}^{D+1}$$

$$\tilde{\mathbf{x}} = \begin{bmatrix} 1 \\ \mathbf{x} \end{bmatrix}$$

$$\in \mathbb{R}^{D+1}$$

$$\tilde{\mathbf{W}} = [\tilde{\mathbf{w}}_1 \dots \tilde{\mathbf{w}}_K]$$

$$\in \mathbb{R}^{(D+1) \times K}$$

<https://powcoder.com>

- we get for the (vector valued) discriminant function

$f(\mathbf{x}) = \tilde{\mathbf{W}}^T \tilde{\mathbf{x}} \in \mathbb{R}^K$

(if it helps, you may think of $\tilde{\mathbf{W}}^T$ as a **vector-valued** function).

- For a new input \mathbf{x} , the class is then defined by the index of the largest value in the row vector $\mathbf{y}(\mathbf{x})$



Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

- Given a training set $\{\mathbf{x}_n, \mathbf{t}_n\}$ where $n = 1, \dots, N$, and \mathbf{t}_n is the class in the 1-of-K coding scheme.

- Define a matrix \mathbf{T} where row n corresponds to \mathbf{t}_n^\top .

- The sum-of-squares error can now be written as
(check that $\text{tr}\{\mathbf{A}^\top \mathbf{A}\} = \|\mathbf{A}\|_F^2$)

$$E_D(\tilde{\mathbf{W}}) = \frac{1}{2} \text{tr} \left\{ (\tilde{\mathbf{X}}\tilde{\mathbf{W}} - \mathbf{T})^\top (\tilde{\mathbf{X}}\tilde{\mathbf{W}} - \mathbf{T}) \right\}$$

- The minimum of $E_D(\tilde{\mathbf{W}})$ will be reached for

$$\tilde{\mathbf{W}} = (\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}})^{-1} \tilde{\mathbf{X}}^\top \mathbf{T} = \tilde{\mathbf{X}}^\dagger \mathbf{T}$$

where $\tilde{\mathbf{X}}^\dagger$ is the pseudo-inverse of $\tilde{\mathbf{X}}$.



Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

- The discriminant function $\mathbf{y}(\mathbf{x})$ is therefore

Assignment Project Exam Help

where $\tilde{\mathbf{X}}$ is given by the training data, and $\tilde{\mathbf{x}}$ is the new input.

- Interesting property: If for every \mathbf{t}_n the same linear constraint $\mathbf{a}^\top \mathbf{t}_n + b = 0$ holds, then the prediction $\mathbf{y}(\mathbf{x})$ will also obey the same constraint

Add WeChat powcoder

- For the 1-of- K coding scheme, the sum of all components in \mathbf{t}_n is one, and therefore all components of $\mathbf{y}(\mathbf{x})$ will sum to one. BUT: the components are not probabilities, as they are not constraint to the interval $(0, 1)$.

Deficiencies of the Least Squares Approach

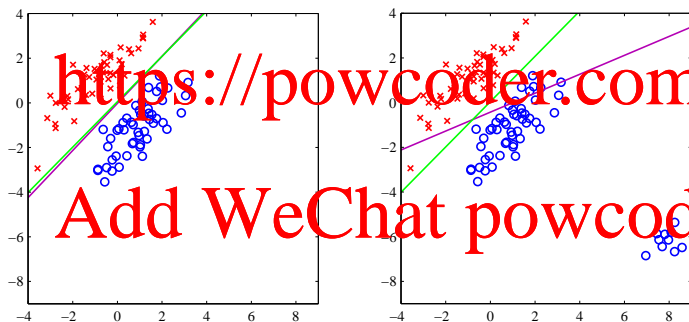


Magenta curve :

Decision Boundary for the least squares approach

Green curve :

Decision boundary for the logistic regression (described later)



(Imagine heat-maps of the quadratic penalty function, similarly to those of the linear functions earlier in the slides.)

Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

Deficiencies of the Least Squares Approach



Left plot : Decision Boundary for least squares
Right plot : Boundaries for logistic regression (described later)



Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm





Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

Assignment Project Exam Help

- View linear classification as dimensionality reduction.

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$$

If $y \geq -w_0$ then class \mathcal{C}_1 , otherwise \mathcal{C}_2 .

- But there are many projections from a 2-dimensional input space onto one dimension.
- Projection always means loss of information.
- For classification we want to preserve the class separation in one dimension.
- Can we find a projection which maximally preserves the class separation ?

<https://powcoder.com>
Add WeChat powcoder

Fisher's Linear Discriminant

Statistical Machine
Learning

© 2020

Ong & Walder & Webers
Data61 \ CSIRO
The Australian National
University



Classification

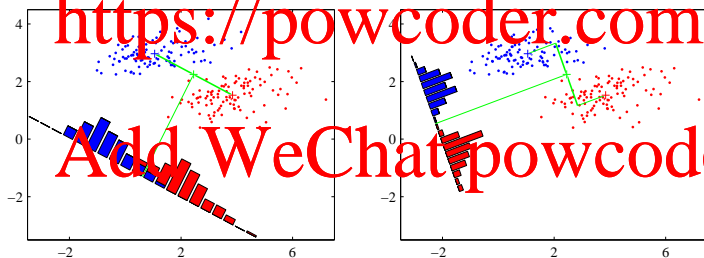
Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

Samples from two classes in a two-dimensional input space
and their histogram when projected to two different
one-dimensional spaces.



<https://powcoder.com>

Add WeChat powcoder

Fisher's Linear Discriminant - First Try

Statistical Machine
Learning

© 2020

Ong & Walder & Webers
Data61 \ CSIRO
The Australian National
University



Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm


- Given N_1 input data of class C_1 , and N_2 input data of class C_2 , calculate the centres of the two classes

$$\mathbf{m}_1 = \frac{1}{N_1} \sum_{n \in C_1} \mathbf{x}_n, \quad \mathbf{m}_2 = \frac{1}{N_2} \sum_{n \in C_2} \mathbf{x}_n$$

- Choose \mathbf{w} so as to maximise the separation of the projected class means

$$\mathbf{m}_1 - \mathbf{m}_2 = \mathbf{w}^T (\mathbf{m}_1 - \mathbf{m}_2)$$

- Problem with non-uniform covariance



Add WeChat powcoder




- Measure also the within-class variance for each class

$$s_k^2 = \sum_{y_n \in \mathcal{C}_k} (y_n - m_k)^2$$

where $y_n = \mathbf{w}^\top \mathbf{x}_n$.

- Maximise the Fisher criterion

$$J(\mathbf{w}) = \frac{(\bar{y}_1 - \bar{y}_2)^2}{s_1^2 + s_2^2}$$



Add WeChat powcoder

Primer: Bilinear form with a Covariance Matrix

Let

$$\boldsymbol{\mu} = \mathbb{E}[\mathbf{x}]$$

$$\Sigma = \text{cov}[\mathbf{x}]$$

$$= \mathbb{E}[(\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^T].$$

Then

$$\text{var}[\mathbf{w}^T \mathbf{x}] = \mathbb{E}[(\mathbf{w}^T \mathbf{x} - \mathbb{E}[\mathbf{w}^T \mathbf{x}])^2]$$

$$= \mathbb{E}[(\mathbf{w}^T \mathbf{x} - \mathbf{w}^T \mathbb{E}[\mathbf{x}])^2]$$

$$= \mathbb{E}[(\mathbf{w}^T \mathbf{x} - \mathbf{w}^T \boldsymbol{\mu})^2]$$

$$= \mathbb{E}[(\mathbf{w}^T \mathbf{x} - \mathbf{w}^T \boldsymbol{\mu})(\mathbf{w}^T \mathbf{x} - \mathbf{w}^T \boldsymbol{\mu})]$$

$$= \mathbb{E}[(\mathbf{w}^T \mathbf{x} - \mathbf{w}^T \boldsymbol{\mu})(\mathbf{x}^T \mathbf{w} - \boldsymbol{\mu}^T \mathbf{w})]$$

$$= \mathbb{E}[\mathbf{w}^T (\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^T \mathbf{w}]$$

$$= \mathbf{w}^T \mathbb{E}[(\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^T] \mathbf{w}$$

$$= \mathbf{w}^T \Sigma \mathbf{w}.$$





Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

- The Fisher criterion can be rewritten as

$$J(\mathbf{w}) = \frac{\mathbf{w}^T \mathbf{S}_B \mathbf{w}}{\mathbf{w}^T \mathbf{S}_W \mathbf{w}}$$

- \mathbf{S}_B is the **between-class** covariance

$$\mathbf{S}_B = (\mathbf{m}_2 - \mathbf{m}_1)(\mathbf{m}_2 - \mathbf{m}_1)^T$$

so by the previous slide, the numerator of $J(\mathbf{w})$ is:
the variance of the projection of the means

- \mathbf{S}_W is the **within-class** covariance

$$\mathbf{S}_W = \sum_{n \in \mathcal{C}_1} (\mathbf{x}_n - \mathbf{m}_1)(\mathbf{x}_n - \mathbf{m}_1)^T + \sum_{n \in \mathcal{C}_2} (\mathbf{x}_n - \mathbf{m}_2)(\mathbf{x}_n - \mathbf{m}_2)^T$$

so so by the previous slide and

$\mathbf{w}^T (A + B) \mathbf{w} = \mathbf{w}^T A \mathbf{w} + \mathbf{w}^T B \mathbf{w}$, the denominator of $J(\mathbf{w})$ is:
(the variance of the projection of the points in class \mathcal{C}_1) +
(the variance of the projection of the points in class \mathcal{C}_2)



Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

Assignment Project Exam Help

- The Fisher criterion

$$J(\mathbf{w}) = \frac{\mathbf{w}^\top \mathbf{S}_B \mathbf{w}}{\mathbf{w}^\top \mathbf{S}_W \mathbf{w}}$$

has a maximum for Fisher's linear discriminant

<https://powcoder.com>

$$\mathbf{w} \propto \mathbf{S}_W^{-1}(\mathbf{m}_2 - \mathbf{m}_1)$$

- Fisher's linear discriminant is NOT a discriminant, but can be used to construct one by choosing a threshold y_0 in the projection space.

Add WeChat powcoder

Fisher's Discriminant For Multi-Class



- Assume that the dimensionality of the input space D is greater than the number of classes K .
- Use $D' > 1$ linear 'features' $y_k = \mathbf{w}_k^\top \mathbf{x}$ and write everything in vector form (with no bias term)

$$\mathbf{y} = \mathbf{W}^\top \mathbf{x}.$$

- The within-class covariance is then the sum of the covariances for all K classes

$$\mathbf{S}_W = \sum_{k=1}^K \mathbf{S}_k$$

where

$$\mathbf{S}_k = \sum_{n \in \mathcal{C}_k} (\mathbf{x}_n - \mathbf{m}_k)(\mathbf{x}_n - \mathbf{m}_k)^\top$$
$$\mathbf{m}_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{x}_n$$



- Between-class covariance

$$S_B = \sum_{k=1}^K N_k (\mathbf{m}_k - \mathbf{m})(\mathbf{m}_k - \mathbf{m})^T$$

where \mathbf{m} is the total mean of the input data

$$\mathbf{m} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n.$$

- One possible way to define a function of \mathbf{W} which is large when the between-class covariance is large and the within-class covariance is small is given by

$$J(\mathbf{W}) = \text{tr} \{ (\mathbf{W}^T \mathbf{S}_W \mathbf{W})^{-1} (\mathbf{W}^T \mathbf{S}_B \mathbf{W}) \}$$

- The maximum of $J(\mathbf{W})$ is determined by the D' eigenvectors of $\mathbf{S}_W^{-1} \mathbf{S}_B$ with the largest eigenvalues.

The Perceptron Algorithm

Statistical Machine
Learning

© 2020

Ong & Walder & Webers
Data61 \ CSIRO
The Australian National
University



Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

- Frank Rosenblatt (1928 - 1969)

- "Principles of neurodynamics: Perceptions and the theory of brain mechanisms" (Spartan Books, 1962)



<https://powcoder.com>

Add WeChat powcoder

The Perceptron Algorithm

Statistical Machine
Learning

© 2020
Ong & Walder & Webers
Data61 \ CSIRO
The Australian National
University



Assignment Project Exam Help

- Perceptron ("MARK 1") was the first computer which could learn new skills by trial and error

Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm



<https://powcoder.com>

Add WeChat powcoder

The Perceptron Algorithm



Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

- Two class model
- Create feature vector $\phi(\mathbf{x})$ by a fixed nonlinear transformation of the input \mathbf{x} .
- Generalised linear model

$$y(\mathbf{x}) = f(\mathbf{w}^\top \phi(\mathbf{x}))$$

with $\phi(\mathbf{x})$ containing some bias element $\phi_0(\mathbf{x}) = 1$

- nonlinear **activation** function

$$f(a) = \begin{cases} +1, & a \geq 0 \\ -1, & a < 0 \end{cases}$$

- Target coding for perceptron

$$t = \begin{cases} +1, & \text{if } C_1 \\ -1, & \text{if } C_2 \end{cases}$$



Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

Assignment Project Exam Help

- Idea: Minimise total number of misclassified patterns.
- Problem: As a function of \mathbf{w} , this is piecewise constant and therefore the gradient is zero almost everywhere.
- Better idea: Using the $(-1, +1)$ target coding scheme, we want all patterns to satisfy $\mathbf{w}^T \phi(\mathbf{x}_n) t_n > 0$.
- **Perceptron Criterion**: Add the errors for all patterns belonging to the set of misclassified patterns \mathcal{M}

Add WeChat powcoder

$$E_P(\mathbf{w}) = - \sum_{n \in \mathcal{M}} \mathbf{w}^T \phi(\mathbf{x}_n) t_n$$



- Perceptron Criterion (with notation $\phi_n = \phi(\mathbf{x}_n)$)

$$E_P(\mathbf{w}) = \sum_{n \in \mathcal{M}} y_n \phi_n^T \mathbf{w} t_n \\ \equiv E_P^{(n)}(\mathbf{w})$$

- One iteration at step τ

- 1 Choose a training data point index n
(uniformly at random or by cycling through the data)
- 2 Update the weight vector \mathbf{w} by

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla E_P^{(n)}(\mathbf{w})$$

where

$$\nabla E_P^{(n)}(\mathbf{w}) = \begin{cases} -\phi_n t_n & \text{if } (\mathbf{w}^{(\tau)})^\top \phi(\mathbf{x}_n) \cdot t_n \leq 0 \\ 0 & \text{otherwise.} \end{cases}$$

- As $y(\mathbf{x}, \mathbf{w})$ is invariant to the norm of \mathbf{w} , we may set $\eta = 1$.

The Perceptron Algorithm - Update 1

Statistical Machine
Learning

© 2020
Ong & Walder & Webers
Data61 \ CSIRO
The Australian National
University



Classification

Generalised Linear
Model

Discriminant Functions

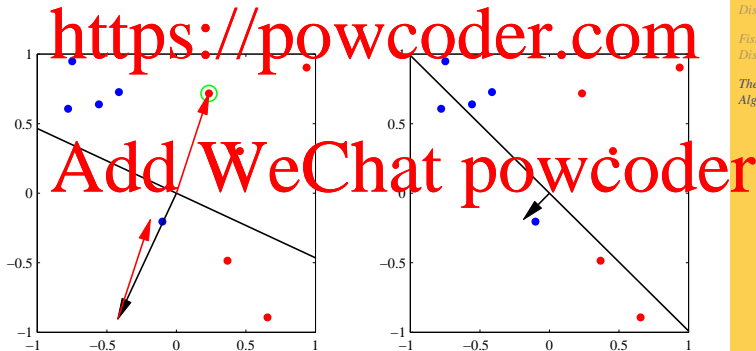
Fisher's Linear
Discriminant

The Perceptron
Algorithm

Update of the perceptron weights from a misclassified pattern

Assignment Project Exam Help

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} + \phi_n t_n$$



The Perceptron Algorithm - Update 2

Statistical Machine
Learning

© 2020

Ong & Walder & Webers
Data61 \ CSIRO
The Australian National
University



Classification

Generalised Linear
Model

Discriminant Functions

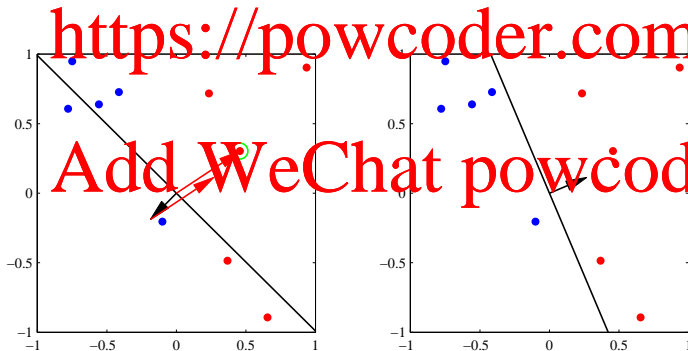
Fisher's Linear
Discriminant

The Perceptron
Algorithm

Update of the perceptron weights from a misclassified pattern

Assignment Project Exam Help

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} + \phi_n t_n$$



The Perceptron Algorithm - Convergence



Classification

Generalised Linear
Model

Discriminant Functions

Fisher's Linear
Discriminant

The Perceptron
Algorithm

- Does the algorithm converge ?

- For a single update step, letting $\eta = 1$, and considering the error from a single point,

$$-\mathbf{w}^{(\tau+1)T} \phi_n t_n = -\mathbf{w}^{(\tau)T} \phi_n t_n - (\phi_n t_n)^\top \phi_n t_n < -\mathbf{w}^{(\tau)T} \phi_n t_n$$

because $(\phi_n t_n)^\top \phi_n t_n = \|\phi_n t_n\|^2 > 0$. In other words, gradient descent on a linear function decreases that function.

- BUT: contributions to the error from the other misclassified patterns might have increased
- AND: some correctly classified patterns might now be misclassified.
- **Perceptron Convergence Theorem** : If the training set is linearly separable, the perceptron algorithm is guaranteed to find a solution in a finite number of steps.