



Outlines

Overview
Introduction
Linear Algebra
Probability
Linear Regression 1
Linear Regression 2
Linear Classification 1
Linear Classification 2
Kernel Methods
Sparse Kernel Methods
Mixture Models and EM 1
Mixture Models and EM 2
Neural Networks 1
Neural Networks 2
Principal Component Analysis
Autoencoders
Graphical Models 1
Graphical Models 2
Graphical Models 3
Sampling
Sequential Data 1
Sequential Data 2

Statistical Machine Learning

Assignment Project Exam Help

Christian Walder

Machine Learning Research Group

CSIRO Data61

and

College of Engineering and Computer Science

The Australian National University

Canberra

Semester One, 2020.

<https://powcoder.com>

Add WeChat powcoder

(Many figures from C. M. Bishop, "Pattern Recognition and Machine Learning")



Neural Networks

Weight-space Symmetries

Parameter Optimisation

Gradient Descent
Optimisation

Assignment Project Exam Help

Part VII

<https://powcoder.com>

Add WeChat powcoder



Neural Networks

Weight-space Symmetries

Parameter Optimisation

Gradient Descent
Optimisation

Assignment Project Exam Help

<https://powcoder.com>

- Play a crucial role in the algorithms explored so far
- Previously (e.g. Linear Regression and Linear Classification): were fixed before learning starts.
- Now for Neural Networks: number of basis functions fixed, parameters of the basis functions are adaptive
- Later in kernel methods: center basis functions on the data / have an infinite number of effective basis functions (e.g. Support Vector Machines).



Neural Networks

Weight-space Symmetries

Parameter Optimisation

Gradient Descent
Optimisation

Assignment Project Exam Help

- The functional form of the network model (including special parametrisation of the basis functions).
- How to determine the network parameters within the maximum likelihood framework?
(Solution of a nonlinear optimisation problem.)
- **Error backpropagation** : efficiently evaluate the derivatives of the log likelihood function with respect to the network parameters.
- Various approaches to regularise neural networks.

<https://powcoder.com>

Add WeChat powcoder



Neural Networks

Weight-space Symmetries

Parameter Optimisation

Gradient Descent
Optimisation

- Same goal as before: e.g. for regression, decompose

$$t(\mathbf{x}) \equiv y(\mathbf{x}, \mathbf{w}) + \epsilon$$

where ϵ is the noise.

- (Generalised) Linear Model

$$y(\mathbf{x}, \mathbf{w}) = f\left(\sum_{j=0}^M w_j \phi_j(\mathbf{x})\right)$$

where $\phi = (\phi_0, \dots, \phi_M)^T$ is the fixed model basis and $\mathbf{w} = (w_0, \dots, w_M)^T$ are the model parameter.

- For regression: $f(\cdot)$ is the identity function.
- For classification: $f(\cdot)$ is a nonlinear activation function.
- Goal : Let $\phi_j(\mathbf{x})$ depend on parameters, and then adjust these parameters together with \mathbf{w} .

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder


Feed-forward Network Functions



- Goal : Let $\phi_j(\mathbf{x})$ depend on parameters, and then adjust these parameters together with \mathbf{w} .

- Many ways to do this.
- Neural networks use basis functions which follow the same form as the (generalised) linear model.

- EACH basis function is itself a nonlinear function of an adaptive linear combination of the inputs



Add WeChat powcoder



- Construct M linear combinations of the input variables x_1, \dots, x_D in the form


$$a_j = \sum_{i=1}^D \underbrace{w_{ji}^{(1)}}_{\text{weights}} x_i + \underbrace{w_{j0}^{(1)}}_{\text{bias}} \quad j = 1, \dots, M$$

$\underbrace{a_j}_{\text{activations}}$

- Apply a differentiable, nonlinear **activation function** $h(\cdot)$ to get the output of the **hidden units**

$$z_j = h(a_j)$$

- $h(\cdot)$ is typically sigmoid, tanh, or more recently ReLU(x) = max(x , 0).





Neural Networks

Weight-space Symmetries

Parameter Optimisation

Gradient Descent
Optimisation

- Outputs of the hidden units are again linearly combined


$$a_k = \sum_{j=1}^M w_{kj}^{(2)} z_j + w_{k0}^{(2)} \quad k=1 \dots K$$

- Apply again a differentiable, nonlinear **activation function** $g(\cdot)$ to get the network outputs y_k

<https://powcoder.com>

$$y_k = g(a_k)$$

Add WeChat powcoder






- The activation function $g(\cdot)$ is determined by the nature of the data and the distribution of the target variables.
- For standard regression: $g(\cdot)$ is the identity so $y_k = a_k$.
- For multiple binary classification, $g(\cdot)$ is a logistic sigmoid:

$$y_k = \sigma(a_k) = \frac{1}{1 + \exp(-a_k)}$$

- Recall from generative classification model perspective:

$$a_k(\mathbf{x}) = \ln \frac{p(\mathbf{x}, \mathcal{C}_{k1})}{p(\mathbf{x}, \mathcal{C}_{k2})}$$





Neural Networks

Weight-space Symmetries

Parameter Optimisation


Gradient Descent
Optimisation

- Combine all transformations into one formula

$$y_k(\mathbf{x}, \mathbf{w}) = g \left(\sum_{j=1}^M w_{kj}^{(2)} h \left(\sum_{i=1}^D w_{ji}^{(1)} x_i + w_{j0}^{(1)} \right) + w_{k0}^{(2)} \right)$$

where \mathbf{w} contains all weight and bias parameters.

<https://powcoder.com>



Add WeChat powcoder




- As before, the biases can be absorbed into the weights by introducing an extra input $x_0 = 1$ and a hidden unit $z_0 = 1$.

Assignment Project Exam Help

$$y_k(\mathbf{x}, \mathbf{w}) = g \left(\sum_{j=0}^M w_{kj}^{(2)} h \left(\sum_{i=0}^D w_{ji}^{(1)} x_i \right) \right)$$

where \mathbf{w} now contains all weight and bias parameters

Add WeChat powcoder




Comparison to Perceptron

- A neural network looks like a **multilayer perceptron**.
- But perceptron's nonlinear activation function was a step function — neither smooth nor differentiable.

$$f(a) = \begin{cases} +1, & a \geq 0 \\ -1, & a < 0 \end{cases}$$

- The activation functions $f(\cdot)$ and $g(\cdot)$ of a neural network are smooth and differentiable.





Neural Networks

Weight-space Symmetries

Parameter Optimisation

Gradient Descent
Optimisation

Assignment Project Exam Help

- If all activation functions are linear functions then there exists an equivalent network without hidden units.

(Composition of linear functions is a linear function.)

- But if the number of hidden units in this case is smaller than the number of input or output units, the resulting linear function are not the most general.

- Dimensionality reduction.

- c.f. Principal Component Analysis (upcoming lecture).

- Generally, most neural networks use nonlinear activation functions as the goal is to approximate a nonlinear mapping from the input space to the outputs.

<https://powcoder.com>

Add WeChat powcoder

Neural Networks as Universal Function Approximators



Neural Networks

Weight-space Symmetries

Parameter Optimisation

Gradient Descent
Optimisation

Assignment Project Exam Help

- Feed-forward neural networks are **universal approximators**.
- Example: A two-layer neural network with linear outputs can uniformly approximate any continuous function on a compact input domain to arbitrary accuracy if it has enough hidden units.
- Holds for a wide range of hidden unit activation functions.
- Remaining big question: Where do we get the appropriate settings for the weights from? With other words, how do we learn the weights from training examples?

<https://powcoder.com>

Add WeChat powcoder



Neural Networks

Weight-space Symmetries

Parameter Optimisation


Gradient Descent
Optimisation

- Neural network approximating

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder




Two-layer network with 3 hidden units (tanh activation functions) and linear outputs trained on 50 data points sampled from the interval $(-1, 1)$. Red: resulting output. Dashed: Output of the hidden units.



- Neural network approximating

Assignment Project Exam Help

$$f(x) = \sin(x)$$



<https://powcoder.com>

Add WeChat powcoder

Two-layer network with 3 hidden units (tanh activation functions) and linear outputs trained on 50 data points sampled from the interval $(-1, 1)$. Red: resulting output. Dashed: Output of the hidden units.



Neural Networks

Weight-space Symmetries

Parameter Optimisation


Gradient Descent
Optimisation

- Neural network approximating

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Two-layer network with 3 hidden units (\tanh activation functions) and linear outputs trained on 50 data points sampled from the interval $(-1, 1)$. Red: resulting output. Dashed: Output of the hidden units.


- Neural network approximating Heaviside function

Assignment Project Exam Help

$$f(x) = \begin{cases} 1, & x \geq 0 \\ 0, & x < 0 \end{cases}$$

<https://powcoder.com>

Add WeChat powcoder



Two-layer network with 3 hidden units (tanh activation functions) and linear outputs trained on 50 data points sampled from the interval $(-1, 1)$. Red: resulting output. Dashed: Output of the hidden units.



Neural Networks

Weight-space Symmetries

Parameter Optimisation

Gradient Descent
Optimisation




Neural Networks

Weight-space Symmetries

Parameter Optimisation

Gradient Descent
Optimisation

- Hidden layer nodes represent parametrised basis functions



<https://powcoder.com>

Add WeChat powcoder

$$z = \sigma(w_0 + w_1x_1 + w_2x_2) \text{ for } (w_0, w_1, w_2) = (0.0, 1.0, 0.1)$$




Neural Networks

Weight-space Symmetries

Parameter Optimisation

Gradient Descent
Optimisation

- Hidden layer nodes represent parametrised basis functions



<https://powcoder.com>

Add WeChat powcoder

$$z = \sigma(w_0 + w_1x_1 + w_2x_2) \text{ for } (w_0, w_1, w_2) = (0.0, 0.1, 1.0)$$

Variable Basis Functions in a Neural Networks

Statistical Machine
Learning

© 2020

Ong & Walder & Webers
Data61 \ CSIRO
The Australian National
University




Neural Networks

Weight-space Symmetries

Parameter Optimisation

Gradient Descent
Optimisation

- Hidden layer nodes represent parametrised basis functions



<https://powcoder.com>

Add WeChat powcoder

$$z = \sigma(w_0 + w_1x_1 + w_2x_2) \text{ for } (w_0, w_1, w_2) = (0.0, -0.5, 0.5)$$




Neural Networks

Weight-space Symmetries

Parameter Optimisation

Gradient Descent
Optimisation

- Hidden layer nodes represent parametrised basis functions




<https://powcoder.com>

Add WeChat powcoder

$$z = \sigma(w_0 + w_1x_1 + w_2x_2) \text{ for } (w_0, w_1, w_2) = (10.0, -0.5, 0.5)$$




- Neural network for two-class classification.
- 2 inputs, 2 hidden units with \tanh activation function, 1 output with logistic sigmoid activation function



Red: $y = 0.5$ decision boundary. Dashed blue: $z = 0.5$ hidden unit contours. Green: Optimal decision boundary from the known data distribution.




- Given a set of weights \mathbf{w} . This fixes a mapping from the input space to the output space.
- Does there exist another set of weights realising the same mapping?
- Assume \tanh activation function for the hidden units. As \tanh is an odd function: $\tanh(-a) = -\tanh(a)$.
- Change the sign of all inputs to a hidden unit and outputs of this hidden unit: Mapping stays the same.





- M hidden units, therefore 2^M equivalent weight vectors.
- Furthermore, exchange all of the weights going into and out of a hidden unit with the corresponding weights of another hidden unit. Mapping stays the same. $M!$ symmetries.
- Overall weight space symmetry : $M! 2^M$

M	1	2	3	4	5	6	7
$M! 2^M$	2	8	48	384	3840	46080	645120



Add WeChat powcoder



Neural Networks

Weight-space Symmetries

Parameter Optimisation

Gradient Descent
Optimisation

Assignment Project Exam Help

- Assume the error $E(\mathbf{w})$ is a smooth function of the weights.
- Smallest value will occur at a **critical point** for which

$\nabla E(\mathbf{w}) = 0$.
<https://powcoder.com>

- This could be a minimum, maximum, or saddle point.
- Furthermore, because of symmetry in weight space, there are at least $M \cdot 2^M$ other critical points with the same value for the error.

Add WeChat powcoder




Definition (Global Minimum)

A point \mathbf{w}^* for which the error $E(\mathbf{w}^*)$ is smaller than any other error $E(\mathbf{w})$.

Definition (Local Minimum)

A point \mathbf{w}^* for which the error $E(\mathbf{w}^*)$ is smaller than any other error $E(\mathbf{w})$ in some neighbourhood of \mathbf{w}^* .




Add WeChat powcoder



- Finding the global minimum is difficult in general (would have to check everywhere) unless the error function comes from a special class (e.g. smooth convex functions have only one local minimum).
- Error functions for neural networks are not convex (symmetries!).
- But finding a local minimum might be sufficient.
- Use iterative methods with weight vector update $\Delta \mathbf{w}^{(\tau)}$ to find a local minimum.

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} + \Delta \mathbf{w}^{(\tau)}$$



Add WeChat powcoder



- Around a stationary point \mathbf{w}^* we can approximate

$$E(\mathbf{w}) \simeq E(\mathbf{w}^*) + \frac{1}{2}(\mathbf{w} - \mathbf{w}^*)^T \mathbf{H}(\mathbf{w} - \mathbf{w}^*),$$

where the Hessian \mathbf{H} is evaluated at \mathbf{w}^* so that

$$H_{ij} = \left. \frac{\partial^2}{\partial w_i \partial w_j} E(\mathbf{w}) \right|_{\mathbf{w}=\mathbf{w}^*}$$

- Using a set $\{\mathbf{u}_i\}$ of orthonormal eigenvectors of \mathbf{H} ,

$$\mathbf{H}\mathbf{u}_i = \lambda_i \mathbf{u}_i,$$

to expand

$$\mathbf{w} - \mathbf{w}^* = \sum_i \alpha_i \mathbf{u}_i.$$

- We get

$$E(\mathbf{w}) = E(\mathbf{w}^*) + \frac{1}{2} \sum_i \lambda_i \alpha_i^2.$$




- Around a minimum \mathbf{w}^* we can approximate

$$E(\mathbf{w}) = E(\mathbf{w}^*) + \frac{1}{2} \sum_i \lambda_i \alpha_i^2.$$

<https://powcoder.com>

Add WeChat powcoder





Neural Networks


Weight-space Symmetries

Parameter Optimisation

Gradient Descent
Optimisation

Assignment Project Exam Help

- Around a minimum \mathbf{w}^* , the Hessian \mathbf{H} must be positive definite if evaluated at \mathbf{w}^* .



- This explains why the Laplace approximation always yields a valid covariance matrix.

Gradient Information improves Performances



Neural Networks

Weight-space Symmetries

Parameter Optimisation

Gradient Descent
Optimisation

- Hessian is symmetric and contains $W(W + 1)/2$ independent entries where W is the total number of weights in the network.

- If we use function evaluations only:

- Need to gather this $O(W^2)$ pieces of information by doing $O(W^2)$ number of function evaluations each of which cost $O(W)$ time, for an overall cost of order $O(W^3)$.

- If we use gradients of the function:

- Surprisingly the gradient ∇E also costs only $O(W)$ time, although it provides W pieces of information.
 - We now need only $O(W)$ steps, so the order of time complexity is reduced to $O(W^2)$.

FYI only: In general we have the “cheap gradient principle”. See (Griewank, A., 2000. *Evaluating Derivatives: Principles and Techniques of Algorithmic Differentiation*, Section 5.1).



Neural Networks

Weight-space Symmetries

Parameter Optimisation

Gradient Descent
Optimisation

Assignment Project Exam Help

- Batch processing : Update the weight vector with a small step in the direction of the negative gradient

$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla E(\mathbf{w}^{(\tau)})$
<https://powcoder.com>


where η is the learning rate.

- After each step, re-evaluate the gradient $\nabla E(\mathbf{w}^{(\tau)})$ again.
- Gradient Descent has problems in 'long valleys':

Add WeChat powcoder



- Gradient Descent has problems in 'long valleys'.



(c) 2006 P.A. Simionescu

Example of zig-zag of Gradient Descent Algorithm.

Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder



Neural Networks

Weight-space Symmetries

Parameter Optimisation

Gradient Descent
Optimisation

Assignment Project Exam Help

- Use Conjugate Gradient Descent instead of Gradient Descent to avoid zig-zag behaviour.

- Use Newton method which also calculates the inverse Hessian in each iteration (but inverting the Hessian is usually costly).

- Use Quasi-Newton methods (e.g. BFGS) which also calculates an estimate of the inverse Hessian while iterating.

- Even simpler are *momentum* based strategies.

- Run the algorithm from a set of starting points to find the smallest local minimum.

<https://powcoder.com>

Add WeChat powcoder



Neural Networks

Weight-space Symmetries

Parameter Optimisation

Gradient Descent
Optimisation

Assignment Project Exam Help

- Remaining big problem: Error function is defined over the whole training set. Therefore, need to process the whole training set for each calculation of the gradient $\nabla E(\mathbf{w}^{(\tau)})$.
- If the error function is a sum of errors for each data point

$$E(\mathbf{w}) = \sum_{n=1} E_n(\mathbf{w})$$

we can use on-line gradient descent (also called sequential gradient descent or stochastic gradient descent) to update the weights by one data point at a time

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla E_n(\mathbf{w}^{(\tau)}).$$



Neural Networks

Weight-space Symmetries

Parameter Optimisation

Gradient Descent
Optimisation

- Add more hidden layers (deep learning). To make it work we need many of the following tricks:

- Clever **weight initialisation** to ensure the gradient is flowing through the entire network.

- Some links may additively **skip** over one or several subsequent layer(s).

- Favour **ReLU** over e.g. the sigmoid, to avoid **vanishing gradients**.

- Clever regularisation methods such as **dropout**.

- Specific architectures, not further considered here:

- Parameters may be shared, notably as in **convolutional neural networks** for images.
- A state space model with neural network transitions is a **recurrent neural network**.
- **Attention** mechanisms learn to focus on specific parts of an input.